



# COMPSCI 351

## Fundamentals of Database Systems

More SQL: Complex Queries, Triggers, Views,  
and Schema Modification



# Outline

- ▶ More Complex SQL Retrieval Queries
  - ▶ Additional features allow users to specify more complex retrievals from database:
    - ▶ Nested queries, joined tables, and outer joins (in the FROM clause), aggregate functions, and grouping
  - ▶ Specifying Semantic Constraints as Assertions and Actions as Triggers
  - ▶ Views (Virtual Tables) in SQL
  - ▶ Schema Modification in SQL



# Comparisons Involving NULL and Three-Valued Logic

- ▶ Meanings of NULL
  - ▶ Unknown value
    - ▶ Value exists but is not known, or it is not known whether or not the value exists
  - ▶ Unavailable or withheld value
    - ▶ Value exists but is purposely withheld
  - ▶ Not applicable attribute
    - ▶ The attribute does not apply to this tuple, or is undefined for this tuple
- ▶ SQL does not distinguish among different meanings of NULL.



# Comparisons Involving NULL and Three-Valued Logic

- ▶ When a **NULL** value is involved in a comparison, the result is considered to be **UNKNOWN**.
- ▶ SQL uses a **three-valued logic**:
  - ▶ **TRUE, FALSE, and UNKNOWN (like Maybe)**

**Table 7.1** Logical Connectives in Three-Valued Logic

(a)	<b>AND</b>	TRUE	FALSE	UNKNOWN
	TRUE	TRUE	FALSE	UNKNOWN
	FALSE	FALSE	FALSE	FALSE
	UNKNOWN	UNKNOWN	FALSE	UNKNOWN
(b)	<b>OR</b>	TRUE	FALSE	UNKNOWN
	TRUE	TRUE	TRUE	TRUE
	FALSE	TRUE	FALSE	UNKNOWN
	UNKNOWN	TRUE	UNKNOWN	UNKNOWN
(c)	<b>NOT</b>			
	TRUE	FALSE		
	FALSE	TRUE		
	UNKNOWN	UNKNOWN		



# Comparisons Involving NULL and Three-Valued Logic (cont'd.)

- ▶ SQL allows queries that check whether an attribute value is NULL
- ▶ Rather than using = or <> to compare an attribute value to NULL, SQL uses the comparison operator:
  - ▶ IS NULL or IS NOT NULL
  - ▶ Each individual NULL value is considered to be distinct from every other NULL value

**Query 18.** Retrieve the names of all employees who do not have supervisors.

**Q18:**    **SELECT**      Fname, Lname  
              **FROM**        EMPLOYEE  
              **WHERE**      Super\_ssn **IS NULL;**



# Nested Queries, Tuples, and Set/Multiset Comparisons

---

## ▶ Nested queries

- ▶ Complete select-from-where blocks within WHERE clause of another query

## ▶ Outer query and nested subqueries

## ▶ Comparison operator IN

- ▶ Compares value  $v$  with a set (or multiset) of values  $V$
- ▶ Evaluates to TRUE if  $v$  is one of the elements in  $V$



# Nested Queries

- ▶ Some queries require that existing values in the database be fetched and then used in a comparison condition.
  - ▶ can be formulated using nested queries
  - ▶ the other query is called the outer query.

```
Q4A:   SELECT      DISTINCT Pnumber
          FROM        PROJECT
          WHERE       Pnumber IN
                      ( SELECT      Pnumber
                          FROM        PROJECT, DEPARTMENT, EMPLOYEE
                          WHERE       Dnum=Dnumber AND
                                      Mgr_ssn=Ssn AND Lname='Smith' )
                      OR
                      Pnumber IN
                      ( SELECT      Pno
                          FROM        WORKS_ON, EMPLOYEE
                          WHERE       Essn=Ssn AND Lname='Smith' );
```



# Nested Queries

- ▶ SQL allows the use of tuples of values in comparison by:

- ▶ Place them within parentheses

```
SELECT      DISTINCT Essn
FROM        WORKS_ON
WHERE       (Pno, Hours) IN ( SELECT      Pno, Hours
                           FROM        WORKS_ON
                           WHERE      Essn='123456789' );
```

- ▶ In addition to the IN operator, a number of other comparison operators can be used to compare a single value  $v$  to a set (or multiset) of values  $V$ .

- ▶  $=$  ANY (or  $=$  SOME) operator

- ▶ Returns TRUE if the value  $v$  is equal to some value in the set  $V$ , and is hence equivalent to IN



# Nested Queries

- ▶ Use other comparison operators to compare a single value v
- ▶ Other operators that can be combined with ANY (or SOME): >, >=, <, <=, and <>
- ▶ The keyword ALL can also be combined with each of these operators
- ▶ value must exceed all values from nested query

```
SELECT      Lname, Fname
FROM        EMPLOYEE
WHERE       Salary > ALL ( SELECT      Salary
                           FROM        EMPLOYEE
                           WHERE       Dno=5 );
```



# Nested Queries

- ▶ In the case of several levels of nested queries, we may have possible ambiguity among the attribute names (if attributes of the same name exist).
- ▶ Avoid potential errors and ambiguities
  - ▶ Create tuple variables (aliases) for all tables referenced in SQL query

**Query 16.** Retrieve the name of each employee who has a dependent with the same first name and is the same sex as the employee.

```
Q16:   SELECT      E.Fname, E.Lname
          FROM       EMPLOYEE AS E
          WHERE      E.Ssn IN  ( SELECT      Essn
                                FROM        DEPENDENT AS D
                                WHERE      E.Fname=D.Dependent_name
                                          AND E.Sex=D.Sex );
```



# Correlated Nested Queries

- ▶ Two queries are **correlated**:
  - ▶ A condition in the WHERE clause of a nested query references some attributes of a relation declared in the outer query.
  - ▶ Is evaluated once for each tuple in the outer query
- ▶ Queries that are nested using the = or IN comparison operator can be collapsed into one single block, e.g., Q16 can be written as:

**Q16A:**    **SELECT**                E.Fname, E.Lname  
                 **FROM**                EMPLOYEE **AS** E, DEPENDENT **AS** D  
                 **WHERE**                E.Ssn=D.Essn **AND** E.Sex=D.Sex  
                                    **AND** E.Fname=D.Dependent\_name;



# The EXISTS and UNIQUE Functions in SQL for correlating queries

- ▶ EXIST and UNIQUE are Boolean functions that return TRUE or FALSE, hence can be used in the WHERE clause
- ▶ EXISTS and NOT EXISTS functions
  - ▶ Check whether the result of a correlated nested query is empty or not.
- ▶ SQL function UNIQUE (Q)
  - ▶ Returns TRUE if there are no duplicate tuples in the result of query Q



# USE of EXISTS

- ▶ EXIST (Q) returns TRUE if there is at least one tuple in the results of the nested query Q, and return FALSE otherwise.

**Query 7.** List the names of managers who have at least one dependent.

```
Q7:   SELECT      Fname, Lname
        FROM       EMPLOYEE
        WHERE      EXISTS ( SELECT      *
                            FROM       DEPENDENT
                            WHERE      Ssn = Essn )
                    AND
                    EXISTS ( SELECT      *
                            FROM       DEPARTMENT
                            WHERE      Ssn = Mgr_ssn );
```



# USE OF NOT EXISTS

- ▶ To achieve the “for all” (universal quantifier) effect, we use *double negation* this way in SQL:
- ▶ Query: List first and last name of employees who work on ALL projects controlled by Department 5.

Q3A:    **SELECT**      Fname, Lname  
          **FROM**        EMPLOYEE  
          **WHERE**      **NOT EXISTS** ( ( **SELECT**      Pnumber  
                        **FROM**        PROJECT  
                        **WHERE**      Dnum = 5)  
                        **EXCEPT**     ( **SELECT**      Pno  
                        **FROM**        WORKS\_ON  
                        **WHERE**      Ssn = Essn ) );



# Explicit Sets and Renaming of Attributes in SQL

- ▶ can use explicit set of values in WHERE clause,
  - ▶ e.g., retrieve the SSN of all employees who work on project numbers 1, 2, or 3.

**Q17:**      **SELECT**      **DISTINCT** Essn  
                 **FROM**         **WORKS\_ON**  
                 **WHERE**        Pno **IN** (1, 2, 3);

- ▶ Use qualifier AS followed by desired new name
  - ▶ Rename any attribute that appears in the result of a query

**Q8A:**    **SELECT**      E.Lname **AS** Employee\_name, S.Lname **AS** Supervisor\_name  
                 **FROM**         EMPLOYEE **AS** E, EMPLOYEE **AS** S  
                 **WHERE**        E.Super\_ssn=S.Ssn;



# Specifying Joined Tables in the FROM Clause of SQL

## ▶ **Joined table (or joined relation)**

- ▶ Permits users to specify a table resulting from a join operation in the FROM clause of a query
- ▶ The FROM clause in Q1A
  - ▶ Contains a single joined table, which has all the attributes of the first table EMPLOYEE followed by all the attributes of the second table DEPARTMENT.

**Q1A:**    **SELECT**       Fname, Lname, Address  
            **FROM**         (EMPLOYEE JOIN DEPARTMENT ON Dno=Dnumber)  
            **WHERE**       Dname='Research';



# Different Types of JOINed Tables in SQL

- ▶ Allow users to specify different types of join, such as
  - ▶ NATURAL JOIN
  - ▶ Various types of OUTER JOIN (LEFT, RIGHT, FULL )
- ▶ NATURAL JOIN on two relations R and S
  - ▶ No join condition specified
  - ▶ Is equivalent to an implicit EQUIJOIN condition for each pair of attributes with same name from R and S
  - ▶ Each such pair of attributes is included only once in the resulting relation.



# NATURAL JOIN

- ▶ If the names of the join attributes are not the same in the base relations, it is possible to **rename** attributes so that they match, and then to apply the NATURAL JOIN:

**QIB:**    **SELECT**    Fname, Lname, Address  
              **FROM**      (EMPLOYEE **NATURAL JOIN**  
                          (DEPARTMENT AS DEPT (Dname, Dno, Mssn,  
                          Msdate)))  
              **WHERE**    Dname='Research';

- ▶ The above works with EMPLOYEE.Dno = DEPT.Dno as an implicit join condition.



# INNER and OUTER Joins

---

## ▶ INNER JOIN (**versus** OUTER JOIN)

- ▶ Default type of join in a joined table
- ▶ Tuple is included in the result only if a matching tuple exists in the other relation

## ▶ LEFT OUTER JOIN

- ▶ Every tuple in left table must appear in result
- ▶ If no matching tuple, added with NULL values for attributes of right table

## ▶ RIGHT OUTER JOIN

- ▶ Every tuple in right table must appear in result
- ▶ If no matching tuple, padded with NULL values for attributes of left table



# Example: LEFT OUTER JOIN

## ▶ LEFT OUTER JOIN

**Q8B:**    **SELECT**        E.Lname **AS** Employee\_name,  
                            S.Lname **AS** Supervisor\_name  
 **FROM**                  (EMPLOYEE **AS** E **LEFT OUTER JOIN** EMPLOYEE **AS** S  
                            ON E.Super\_ssn = S.Ssn);

- ▶ In some DBMS, a different syntax was used to specify outer joins by using the comparison operators, e.g.,
  - ▶ In Oracle, '+=' , '=+' and '+=+' symbols are used for left, right and full outer join respectively.

**Q8C:**    **SELECT**        E.Lname, S.Lname  
                            **FROM**                  EMPLOYEE E, EMPLOYEE S  
                            **WHERE**                E.Super\_ssn + = S.Ssn;



# Multiway JOIN in the FROM clause

- ▶ FULL OUTER JOIN – combines results of the LEFT and RIGHT OUTER JOIN
- ▶ Can nest JOIN specifications for a multiway join, e.g.,

**Query 2.** For every project located in ‘Stafford’, list the project number, the controlling department number, and the department manager’s last name, address, and birth date.

Q2:    **SELECT**       Pnumber, Dnum, Lname, Address, Bdate  
          **FROM**         PROJECT, DEPARTMENT, EMPLOYEE  
          **WHERE**        Dnum = Dnumber **AND** Mgr\_ssn = Ssn **AND**  
                        Plocation = ‘Stafford’

Q2A:    **SELECT**       Pnumber, Dnum, Lname, Address, Bdate  
          **FROM**         ((PROJECT **JOIN** DEPARTMENT **ON** Dnum = Dnumber)  
                          **JOIN** EMPLOYEE **ON** Mgr\_ssn = Ssn)  
          **WHERE**        Plocation = ‘Stafford’;



# Aggregate Functions in SQL

- ▶ Used to summarize information from multiple tuples into a single-tuple summary
- ▶ Built-in aggregate functions
  - ▶ COUNT, SUM, MAX, MIN, and AVG
- ▶ **Grouping**
  - ▶ Create subgroups of tuples before summarizing
- ▶ To select entire groups, HAVING clause is used
- ▶ Aggregate functions can be used in the SELECT clause or in a HAVING clause



# Renaming Results of Aggregation

- ▶ Following query returns a single row of computed values from EMPLOYEE table:

**Query 19.** Find the sum of the salaries of all employees, the maximum salary, the minimum salary, and the average salary.

Q19:     **SELECT**     **SUM** (Salary), **MAX** (Salary), **MIN** (Salary), **AVG** (Salary)  
             **FROM**       EMPLOYEE;

- ▶ The result can be presented with new names:

We could use AS to rename the column names in the resulting single-row table; for example, as in Q19A.

Q19A:    **SELECT**     **SUM** (Salary) **AS** Total\_Sal, **MAX** (Salary) **AS** Highest\_Sal,  
              **MIN** (Salary) **AS** Lowest\_Sal, **AVG** (Salary) **AS** Average\_Sal  
             **FROM**       EMPLOYEE;



# Aggregate Functions in SQL

- ▶ NULL values are discarded when aggregate functions are applied to a particular column

**Query 20.** Find the sum of the salaries of all employees of the ‘Research’ department, as well as the maximum salary, the minimum salary, and the average salary in this department.

**Q20:**    **SELECT**      **SUM** (Salary), **MAX** (Salary), **MIN** (Salary), **AVG** (Salary)  
              **FROM**         (**EMPLOYEE JOIN DEPARTMENT ON Dno = Dnumber**)  
              **WHERE**       Dname = ‘Research’;

**Queries 21 and 22.** Retrieve the total number of employees in the company (Q21) and the number of employees in the ‘Research’ department (Q22).

**Q21:**    **SELECT**      **COUNT** (\*)  
              **FROM**         **EMPLOYEE**;

**Q22:**    **SELECT**      **COUNT** (\*)  
              **FROM**         **EMPLOYEE, DEPARTMENT**  
              **WHERE**       **DNO = DNUMBER AND DNAME = ‘Research’**;



# Aggregate Functions in SQL

- ▶ Use the COUNT function to count values in a column (attribute) rather than tuples, e.g.,

**Query 23.** Count the number of distinct salary values in the database.

**Q23:**    **SELECT**      **COUNT (DISTINCT Salary)**  
              **FROM**        **EMPLOYEE;**

- ▶ We can specify a correlated nested query with an aggregate function, e.g., to retrieve the names of all employees who have two or more dependents (Query 5)

**Q5:**    **SELECT**      Lname, Fname  
              **FROM**        EMPLOYEE  
              **WHERE**      ( **SELECT**      **COUNT (\*)**  
                        **FROM**        DEPENDENT  
                        **WHERE**      Ssn = Essn ) >= 2;



# Grouping: The GROUP BY Clause

- ▶ **Partition** the relation into nonoverlapping subsets (or groups) of tuples
  - ▶ Each group (partition) consists of the tuples that have the same value of some attributes, called **grouping attribute(s)**
  - ▶ Apply function to each such group independently to produce summary information about each group
- ▶ **GROUP BY clause**
  - ▶ Specifies grouping attributes
  - ▶ **COUNT (\*)** counts the number of rows in the group



## Examples of GROUP BY

- ▶ The grouping attribute must appear in the SELECT clause

**Q24:**      **SELECT**              Dno, **COUNT** (\*), **AVG** (Salary)  
                **FROM**                 EMPLOYEE  
                **GROUP BY**         Dno;

- ▶ If the grouping attribute has NULL as a possible value, then a separate group is created for the null value (e.g., null Dno in the above query).
- ▶ GROUP BY may be applied to the result of a JOIN:

**Q25:**      **SELECT**              Pnumber, Pname, **COUNT** (\*)  
                **FROM**                 PROJECT, WORKS\_ON  
                **WHERE**               Pnumber=Pno  
                **GROUP BY**         Pnumber, Pname;



# Grouping: The GROUP BY and HAVING Clauses

- ▶ **HAVING clause**
  - ▶ Provides a condition to select or reject an entire group
- ▶ **Query 26.** For each project *on which more than two employees work*, retrieve the project number, the project name, and the number of employees who work on the project.

**Q26:**      **SELECT**      Pnumber, Pname, **COUNT (\*)**  
                 **FROM**        PROJECT,WORKS\_ON  
                 **WHERE**      Pnumber=Pno  
                 **GROUP BY** Pnumber, Pname  
                 **HAVING**     **COUNT (\*) > 2;**



# Combining the WHERE and the HAVING Clause

## ▶ Consider the query:

- ▶ we want to count the *total* number of employees whose salaries exceed \$40,000 in each department, but only for departments where more than five employees work.

### ▶ **INCORRECT** QUERY:

```
SELECT      Dno, COUNT (*)
FROM        EMPLOYEE
WHERE       Salary>40000
GROUP BY    Dno
HAVING     COUNT (*) > 5;
```

- ▶ WHERE clause is executed first to select individual or joined tuples
- ▶ HAVING clause is applied later to select individual groups of tuples.



# Combining the WHERE and the HAVING Clause

## Correct Specification of the Query:

- ▶ Note: the WHERE clause applies tuple by tuple whereas HAVING applies to entire group of tuples

**Query 28.** For each department that has more than five employees, retrieve the department number and the number of its employees who are making more than \$40,000.

Q28:

<b>SELECT</b>	Dno, <b>COUNT(*)</b>
<b>FROM</b>	EMPLOYEE
<b>WHERE</b>	Salary > 40000 <b>AND</b> Dno <b>IN</b>
	( <b>SELECT</b> Dno
	<b>FROM</b> EMPLOYEE
	<b>GROUP BY</b> Dno
	<b>HAVING</b> <b>COUNT(*) &gt; 5</b>
<b>GROUP BY</b>	Dno;



## Other SQL Construct: WITH

- ▶ The WITH clause allows a user to define a table that will only be used in a particular query (*may not be available in all SQL based DBMS implementations*)
- ▶ Used for convenience to create a temporary “View” and use that immediately in a query
- ▶ Allows a more straightforward way of looking a step-by-step query
- ▶ This temporary “View” (or table) will be discarded after the query is executed.



# Example of WITH

- ▶ See an alternate approach to performing Q28:

**Query 28.** For each department that has more than five employees, retrieve the department number and the number of its employees who are making more than \$40,000.

```
Q28':    WITH      BIGDEPTS (Dno) AS
              ( SELECT      Dno
                FROM        EMPLOYEE
                GROUP BY    Dno
                HAVING      COUNT (*) > 5)
              SELECT      Dno, COUNT (*)
              FROM        EMPLOYEE
              WHERE       Salary>40000 AND Dno IN BIGDEPTS
              GROUP BY    Dno;
```

- ▶ We define in the WITH clause a temporary table BIGDEPTS, and use this table in the subsequent query.



# Other SQL Construct: CASE

- ▶ SQL also has a CASE construct
  - ▶ Used when a value can be different based on certain conditions.
  - ▶ Can be used in any part of an SQL query where a value is expected
  - ▶ Applicable when querying, inserting or updating tuples
  - ▶ For example, employees are receiving different raises in different departments (A variation of the update U6)

**U6':**      **UPDATE**      **EMPLOYEE**  
                 **SET**               **Salary =**  
                 **CASE**             **WHEN**           **Dno = 5**      **THEN** **Salary + 2000**  
                                **WHEN**           **Dno = 4**      **THEN** **Salary + 1500**  
                                **WHEN**           **Dno = 1**      **THEN** **Salary + 3000**  
                 **ELSE**               **Salary + 0 ;**



# Recursive Queries in SQL

- ▶ An example of a **recursive relationship** between tuples of the same type is the relationship between an employee and a supervisor.
- ▶ This relationship is described by the foreign key `Super_ssn` of the `EMPLOYEE` relation
- ▶ A example of a **recursive operation** is to retrieve all supervisees of a supervisory employee  $e$  at all levels, i.e.,
  - ▶ all employees  $e'$  directly supervised by  $e$ , all employees  $e''$  directly supervised by each employee  $e'$ , all employees  $e'''$  directly supervised by each employee  $e''$ , and so on. Thus the CEO would have each employee in the company as a supervisee in the resulting table. Example shows such table `SUP_EMP` with 2 columns (`Supervisor`, `Supervisee (any level)`):



# An EXAMPLE of RECURSIVE Query

▶ **Q29: WITH RECURSIVE**

```
( SELECT          SUP_EMP (SupSsn, EmpSsn) AS
  FROM           SupervisorSsn, Ssn
  EMPLOYEE
  UNION
  SELECT          S.EmpSsn, E.Ssn
  FROM           EMPLOYEE AS E, SUP_EMP AS S
  WHERE          E.SupervisorSsn = S.EmpSsn )
  SELECT          *
  FROM           SUP_EMP;
```

- ▶ The above query starts with an empty SUP\_EMP and successively builds SUP\_EMP table by computing immediate supervisees first, then second level supervisees, etc. until a **fixed point** is reached and no more supervisees can be added.



# Summary of Expanded Block Structure of SQL Queries

```
SELECT <attribute and function list>
FROM <table list>
[ WHERE <condition> ]
[ GROUP BY <grouping attribute(s)> ]
[ HAVING <group condition> ]
[ ORDER BY <attribute list> ];
```

- ▶ **SELECT clause** lists the attributes or functions to be retrieved
- ▶ **FROM clause** specifies all the tables needed
- ▶ **WHERE clause** specifies the conditions for selection
- ▶ **GRROUP BY** specifies grouping attributes
- ▶ **HAVING** specifies a condition on the grouping
- ▶ **ORDER BY** specifies an order for displaying the result



# Specifying Constraints as Assertions and Actions as Triggers

## ▶ **CREATE ASSERTION**

- ▶ Specify additional types of constraints that are outside scope of the built-in relational model constraints (i.e., *primary and unique keys, entity integrity, and referential integrity*).

## ▶ **CREATE TRIGGER**

- ▶ Specify automatic actions that database system will perform when certain events and conditions occur
- ▶ This type of functionality is generally referred as ***active databases***.



# Specifying General Constraints as Assertions in SQL

## ▶ CREATE ASSERTION

- ▶ Specify a query that selects any tuples that violate the desired condition
- ▶ Use only in cases where it goes beyond a simple CHECK which applies to individual attributes and domains
- ▶ E.g., a constraint specifies that the salary of an employee must not be greater than that of the manager in the same department

```
CREATE ASSERTION SALARY_CONSTRAINT
CHECK ( NOT EXISTS ( SELECT      *
                      FROM        EMPLOYEE E, EMPLOYEE M,
                                  DEPARTMENT D
                     WHERE      E.Salary>M.Salary
                               AND E.Dno=D.Dnumber
                               AND D.Mgr_ssn=M.Ssn ) );
```



# Introduction to Triggers in SQL

- ▶ CREATE TRIGGER statement
  - ▶ Specify the type of action to be taken when certain events occur and when certain conditions are satisfied.
    - ▶ E.g., a manager may want to be informed if an employee's travel expenses exceed a certain limit in the database state
  - ▶ Used to monitor the database, and implement such actions in DBMS
- ▶ Typical trigger has three components which make it a rule for an “active database” :
  - ▶ **Event(s)**: usually database update operations
  - ▶ **Condition**: determines whether rule actions should be executed
  - ▶ **Action**: to be taken, usually SQL statements or transactions, etc.



# Example use of TRIGGERS

- ▶ Suppose we want to check whenever an employee's salary is greater than the salary of his or her direct supervisor in the COMPANY database.
  - ▶ Several events can trigger this rule: inserting a new employee record, changing an employee's salary, or changing an employee's supervisor.
  - ▶ Suppose that the action to take would be to call an external stored procedure SALARY\_VIOLATION,5 which will notify the supervisor.

R5: **CREATE TRIGGER SALARY\_VIOLATION**  
**BEFORE INSERT OR UPDATE OF SALARY, SUPERVISOR\_SSN**  
**ON EMPLOYEE**  
**FOR EACH ROW**  
**WHEN ( NEW.SALARY > ( SELECT SALARY FROM EMPLOYEE**  
**WHERE SSN = NEW.SUPERVISOR\_SSN ) )**  
**INFORM\_SUPERVISOR(NEW.Supervisor\_ssN,**  
**NEW.Ssn );**



# Views (Virtual Tables) in SQL

- ▶ Concept of a view in SQL
  - ▶ Single table derived from other tables (can be base tables or previously defined views)
  - ▶ Does not necessarily exist in physical form, is considered to be a **virtual table** (in contrast to base tables)
  - ▶ As a way of specifying a table that we need to reference frequently, even though it may not exist physically.
- ▶ Once a View is defined, SQL queries can use the View relation in the FROM clause
- ▶ View is always up-to-date
  - ▶ Responsibility of the DBMS and not the user



# Specification of Views in SQL

## ▶ CREATE VIEW command

- ▶ Give table name, list of attribute names, and a query to specify the contents of the view
- ▶ In V1, attributes retain the names from base tables. In V2, attributes are assigned names

V1:	<b>CREATE VIEW</b>	WORKS_ON1
	<b>AS SELECT</b>	Fname, Lname, Pname, Hours
	<b>FROM</b>	EMPLOYEE, PROJECT, WORKS_ON
	<b>WHERE</b>	Ssn = Essn <b>AND</b> Pno = Pnumber;
V2:	<b>CREATE VIEW</b>	DEPT_INFO(Dept_name, No_of_emps, Total_sal)
	<b>AS SELECT</b>	Dname, <b>COUNT (*)</b> , <b>SUM</b> (Salary)
	<b>FROM</b>	DEPARTMENT, EMPLOYEE
	<b>WHERE</b>	Dnumber = Dno
	<b>GROUP BY</b>	Dname;

## ▶ DROP VIEW command

- ▶ Dispose of a view



# View Implementation, View Update, and Inline Views

---

- ▶ Complex problem of efficiently implementing a view for querying
- ▶ **Strategy I: Query modification** approach
  - ▶ Compute the view as and when needed. Do not store permanently
  - ▶ Modify view query into a query on underlying base tables
  - ▶ Disadvantage: inefficient for views defined via complex queries that are time-consuming to execute



# View Materialization

## ▶ **Strategy 2: View materialization**

- ▶ Physically create a temporary view table when the view is first queried or created
- ▶ Keep that table on the assumption that other queries on the view will follow
- ▶ Requires efficient strategy for automatically updating the view table when the base tables are updated

## ▶ **Incremental update strategy for materialized views**

- ▶ DBMS determines what new tuples must be inserted, deleted, or modified in a materialized view table



# View Materialization

---

- ▶ Multiple ways to handle materialization:
  - ▶ **immediate update** strategy updates a view as soon as the base tables are changed
  - ▶ **lazy update** strategy updates the view when needed by a view query
  - ▶ **periodic update** strategy updates the view periodically (in the latter strategy, a view query may get a result that is not up-to-date). This is commonly used in Banks, Retail store operations, etc.



# View Update

- ▶ Update on a view defined on a single table without any aggregate functions
  - ▶ Can be mapped to an update on underlying base table - possible if the primary key is preserved in the view
- ▶ Update not permitted on aggregate views, e.g.,

<b>UV2: UPDATE</b>	<b>DEPT_INFO</b>
<b>SET</b>	Total_sal=100000
<b>WHERE</b>	Dname='Research';
- ▶ cannot be processed because **Total\_sal** is a computed value in the view definition



# View Update and Inline Views

---

- ▶ View involving joins
  - ▶ Often not possible for DBMS to determine which of the updates is intended
- ▶ Clause **WITH CHECK OPTION**
  - ▶ Must be added at the end of the view definition if a view is to be updated to make sure that tuples being updated stay in the view
- ▶ In-line view
  - ▶ Defined in the FROM clause of an SQL query (e.g., we saw its used in the WITH example)



# Views as authorization mechanism

- ▶ SQL query authorization statements (**GRANT** and **REVOKE**) are described in later topics.
- ▶ Views can be used to hide certain attributes or tuples from unauthorized users
- ▶ E.g., For a user who is only allowed to see employee information for those who work for department 5, he or she may only access the view DEPT5EMP:

```
CREATE VIEW DEPT5EMP AS
SELECT *
FROM EMPLOYEE
WHERE Dno = 5;
```



# Schema Change Statements in SQL

## ▶ **Schema evolution commands**

- ▶ Can be used to alter a schema by adding or dropping tables, attributes, constraints and other schema elements.
- ▶ DBA may want to change the schema while the database is operational
- ▶ Does not require recompilation of the database schema
- ▶ Certain checks must be performed by the DBMS to ensure that the changes do not affect the rest of the database and cause inconsistencies.



# The DROP Command

---

- ▶ **DROP command**
  - ▶ Used to drop named schema elements, such as tables, domains, or constraint
- ▶ **Drop behavior options:**
  - ▶ CASCADE and RESTRICT
- ▶ **Example:**
  - ▶ `DROP SCHEMA COMPANY CASCADE;`
  - ▶ This removes the schema and all its elements including tables, views, constraints, etc.
  - ▶ `DROP TABLE DEPENDENT RESTRICT;`
  - ▶ Table is dropped only if not referenced in any constraints



# The ALTER table command

- ▶ **Alter table actions** include:
  - ▶ Adding or dropping a column (attribute)
  - ▶ Changing a column definition
  - ▶ Adding or dropping table constraints
- ▶ **Example:**
  - ▶ `ALTER TABLE COMPANY.EMPLOYEE ADD COLUMN Job VARCHAR(12);`
  - ▶ We must still enter a value for the new attribute `Job` for each individual `EMPLOYEE` tuple.
  - ▶ By using the default clause or the `UPDATE` command



# Dropping Columns, Default Values

- ▶ Drop a column must choose either CASCADE or RESTRICT
  - ▶ CASCADE: all constraints and views that reference the column are dropped automatically
  - ▶ RESTRICT: the drop is successful only when no views or constraints (or other schema elements) reference the column

```
ALTER TABLE COMPANY.EMPLOYEE DROP COLUMN Address CASCADE;
```

- ▶ Default values can be dropped and altered:

```
ALTER TABLE COMPANY.DEPARTMENT ALTER COLUMN Mgr_ssn  
DROP DEFAULT;
```

```
ALTER TABLE COMPANY.DEPARTMENT ALTER COLUMN Mgr_ssn  
SET DEFAULT '333445555';
```



# Adding and Dropping Constraints

- ▶ We can change the constraints specified on a table by
  - ▶ Add or drop a named constraint
  - ▶ For example, to drop the constraint named EMPSUPERFK from the EMPLOYEE table

```
ALTER TABLE COMPANY.EMPLOYEE  
DROP CONSTRAINT EMPSUPERFK CASCADE;
```

- ▶ After removing the constraint, we can redefine a replacement constraint using the ADD CONSTRAINT keyword in the ALTER TABLE statement, e.g.,

```
ALTER TABLE COMPANY.EMPLOYEE  
ADD CONSTRAINT EMPSUPERFK  
FOREIGN KEY(Super_ssn) REFERENCES EMPLOYEE(Ssn)  
ON DELETE SET NULL ON UPDATE CASCADE;
```



# Summary of SQL Syntax

**Table 7.2** Summary of SQL Syntax

---

```
CREATE TABLE <table name> ( <column name><column type> [ <attribute constraint> ]
    { , <column name><column type> [ <attribute constraint> ] }
    [ <table constraint> { , <table constraint> } ] )
```

---

```
DROP TABLE <table name>
```

```
ALTER TABLE <table name> ADD <column name><column type>
```

```
SELECT [ DISTINCT ] <attribute list>
```

```
FROM ( <table name> { <alias> } | <joined table> ) { , ( <table name> { <alias> } | <joined table> ) }
[ WHERE <condition> ]
```

```
[ GROUP BY <grouping attributes> [ HAVING <group selection condition> ] ]
```

```
[ ORDER BY <column name> [ <order> ] { , <column name> [ <order> ] } ]
```

---

```
<attribute list> ::= ( * | ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) )
    { , ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) ) ) ) ) )
```

---

```
<grouping attributes> ::= <column name> { , <column name> }
```

```
<order> ::= ( ASC | DESC )
```

---

```
INSERT INTO <table name> [ ( <column name> { , <column name> } ) ]
( VALUES ( <constant value> , { <constant value> } ) { , ( <constant value> { , <constant value> } ) }
| <select statement> )
```

---

*continued on next slide*



# Summary of SQL Syntax

**Table 7.2** Summary of SQL Syntax

---

DELETE FROM <table name>

[ WHERE <selection condition> ]

---

UPDATE <table name>

SET <column name> = <value expression> { , <column name> = <value expression> }

[ WHERE <selection condition> ]

---

CREATE [ UNIQUE] INDEX <index name>

ON <table name> ( <column name> [ <order> ] { , <column name> [ <order> ] } )

[ CLUSTER ]

---

DROP INDEX <index name>

---

CREATE VIEW <view name> [ ( <column name> { , <column name> } ) ]

AS <select statement>

---

DROP VIEW <view name>

---

NOTE: The commands for creating and dropping indexes are not part of standard SQL.

## ► Notes on the BNF notation:

- <...> for nonterminal symbols, and [...] for optional parts;
- { ... } for repetitions, and (... | ... | ...) for alternatives.



# Summary

---

- ▶ **Complex SQL:**
  - ▶ Nested queries, joined tables (in the FROM clause), outer joins, aggregate functions, grouping
- ▶ **Handling semantic constraints with CREATE ASSERTION and CREATE TRIGGER**
- ▶ **CREATE VIEW statement and materialization strategies**
- ▶ **Schema Modification for the DBAs using ALTER TABLE , ADD and DROP COLUMN, ALTER CONSTRAINT etc .**